



- I delete the early_stopping for now. Because it stopst the training too early like at 16. epoch



- Okk I don't know why but my code start training with too high loss and val acc doesn't increase. So I will start from manos's code and do the changes step by step to see where does the problem come


Starting from the beginning.

- Manos'us code without any change (it stops at 74. Epoch and acc's seems ok)




```
16/16 [=====] - 10s 599ms/step - loss: 0.5544 -  
accuracy: 0.7812 - val_loss: 0.6686 - val_accuracy: 0.7525  
Epoch 73/100  
15/16 [=====>..] - ETA: 0s - loss: 0.6093 - accuracy:  
0.7711  
Epoch 00073: val_loss did not improve from 0.64415  
16/16 [=====] - 9s 594ms/step - loss: 0.5987 -  
accuracy: 0.7792 - val_loss: 0.6675 - val_accuracy: 0.7525  
Epoch 74/100  
15/16 [=====>..] - ETA: 0s - loss: 0.5737 - accuracy:  
0.7578  
Epoch 00074: val_loss did not improve from 0.64415  
16/16 [=====] - 9s 589ms/step - loss: 0.5728 -  
accuracy: 0.7604 - val_loss: 0.6703 - val_accuracy: 0.7475
```

- Manos'us code without any change but with the new dataset (It doesn't go to early stop and actually we see train and val loss still decrease so I will increase the epoch number)



```
Epoch 00099: val_loss did not improve from 0.23544
16/16 [=====] - 13s 823ms/step - loss: 0.5813 -
accuracy: 0.7792 - val_loss: 0.2362 - val_accuracy: 0.9333
Epoch 100/100
15/16 [=====>..] - ETA: 0s - loss: 0.4173 - accuracy:
0.8400
Epoch 00100: val_loss improved from 0.23544 to 0.23248, saving model to
model.weights.best.embedding_network.hdf5
16/16 [=====] - 13s 841ms/step - loss: 0.4082 -
accuracy: 0.8438 - val_loss: 0.2325 - val_accuracy: 0.9333
```

- Manos'us code with input size = 100x100 and epoch = 140



```
Epoch 00135: val_loss did not improve from 0.15591
16/16 [=====] - 12s 774ms/step - loss: 0.3073 - accuracy: 0.9000 - val_loss: 0.1719 - val_accuracy: 0.9472
Epoch 136/140
15/16 [=====>..] - ETA: 0s - loss: 0.2552 - accuracy: 0.8956
Epoch 00136: val_loss did not improve from 0.15591
16/16 [=====] - 12s 772ms/step - loss: 0.2642 - accuracy: 0.8979 - val_loss: 0.1596 - val_accuracy: 0.9528
```



The training stops with early stop at 136 epoch. So 140 epoch was ok we don't need more.

- Changing learning rate from 0.0001 to learning scheduler

```
def lr_scheduler(epoch, lr=0.1):
    i = int(epoch/20)
    if epoch < 20:
        return lr
    else:
        return lr * tf.math.exp(pow(10,-i))



lr_ = tf.keras.callbacks.LearningRateScheduler(lr_scheduler, verbose=1)
```

ULALAAA I did it from scratch but look the accuracy got messed up. I will use some google prepared learning scheduler :D



```
def lr_exp_decay(epoch, lr):
    k = 0.1
    return initial_learning_rate * math.exp(-k*epoch)

lr_ = tf.keras.callbacks.LearningRateScheduler(lr_exp_decay, verbose=1)
```



Okkkk step decay doesn't work for us I don't know whyyy. The possible reason:

In transfer learning, learning rate should be very very low
otherwise the weights change too much even in the first epoch... so we loose the pretrain
benefits.


- Why do we have untrainable params after adding last 3 dense to the vgg network??


```
Total params: 14,714,688  
Trainable params: 14,714,688  
Non-trainable params: 0
```

```
Total params: 15,149,379  
Trainable params: 15,147,843  
Non-trainable params: 1,536
```

In general, all weights are trainable weights. The only built-in layer that has non-trainable weights is the BatchNormalization layer. It uses non-trainable weights to keep track of the mean and variance of its inputs during training. OKK its normal

- **FINAL MODEL**





```

0.0986 - val_accuracy: 0.9694
Epoch 138/140
23/24 [=====>..] - ETA: 0s - loss: 0.1488 - accuracy: 0.9457
Epoch 00138: val_loss improved from 0.09831 to 0.09663, saving model to embedding_network.h5
24/24 [=====] - 11s 471ms/step - loss: 0.1462 - accuracy: 0.9458 - val_loss:
0.0966 - val_accuracy: 0.9694
Epoch 139/140
23/24 [=====>..] - ETA: 0s - loss: 0.2085 - accuracy: 0.9261
Epoch 00139: val_loss improved from 0.09663 to 0.09617, saving model to embedding_network.h5
24/24 [=====] - 11s 473ms/step - loss: 0.2023 - accuracy: 0.9271 - val_loss:
0.0962 - val_accuracy: 0.9722
Epoch 140/140
23/24 [=====>..] - ETA: 0s - loss: 0.2092 - accuracy: 0.9304
Epoch 00140: val_loss did not improve from 0.09617
24/24 [=====] - 11s 463ms/step - loss: 0.2036 - accuracy: 0.9312 - val_loss:
0.0976 - val accuracy: 0.9694

```

```

Test loss: 0.12065759352408349
Test Acc. = 0.9646464646464646

```

pretrain_mano_changes.py
embedding_network.h5

Additional Info I had during the experiments

1)

1 Answer

Active Oldest



In ReduceLROnPlateau, lr changes at the end of previous epoch, In LearningRateScheduler lr changes at the beginning of current epoch.

3



therefore, LearningRateScheduler always wins.



Share Edit Follow Flag

answered Sep 12 '19 at 22:27



Shawn

2)

conda install h5py==2.10.0 otherwise load.model gives "AttributeError: 'str' object has no attribute 'decode' ", while Loading a Keras Saved Model